

ALMOND: Hierarchical Dialogue Policy Learning

Adaptive Language-guided Multi-objective Dialogue with Intent-Drift Monitoring

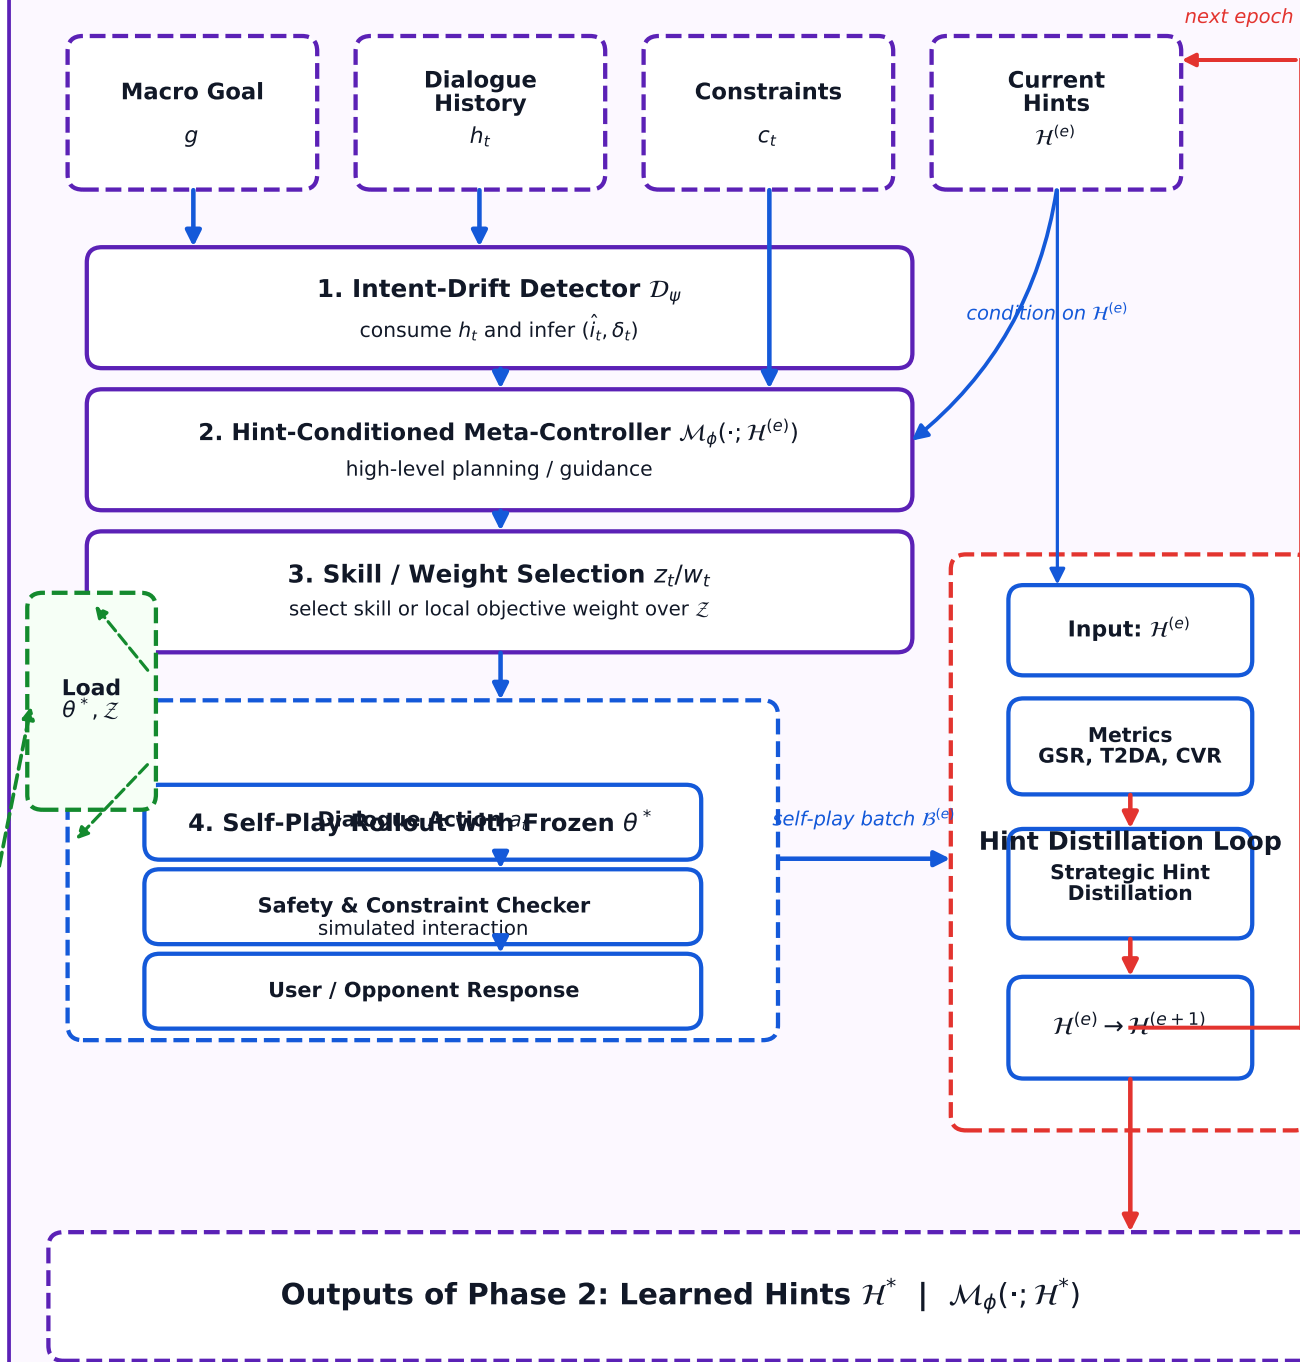
Training Phase 1: Low-Level Skill Learning

learn low-level skill policy only




Training Phase 2: High-Level Meta-Controller Adaptation

freeze low-level policy, run self-play, and update hints only



Inference

use learned components only (no updates)



Legend:



Execution / Data-Control Flow



Learning / Update during training only



Loading / Reuse of learned components



Frozen component